

CS 535 Large Scale Data Analysis

Amit Jain



Big Data, Big Disks, Cheap Computers

- ▶ *"In pioneer days they used oxen for heavy pulling, and when one ox couldn't budge a log, they didn't try to grow a larger ox. We shouldn't be trying for bigger computers, but for more systems of computers."* Rear Admiral Grace Hopper.

Big Data, Big Disks, Cheap Computers

- ▶ *"In pioneer days they used oxen for heavy pulling, and when one ox couldn't budge a log, they didn't try to grow a larger ox. We shouldn't be trying for bigger computers, but for more systems of computers."* Rear Admiral Grace Hopper.
- ▶ *"More data usually beats better algorithms."* Anand Rajaraman.


Big Data, Big Disks, Cheap Computers

- ▶ *"In pioneer days they used oxen for heavy pulling, and when one ox couldn't budge a log, they didn't try to grow a larger ox. We shouldn't be trying for bigger computers, but for more systems of computers."* Rear Admiral Grace Hopper.
- ▶ *"More data usually beats better algorithms."* Anand Rajaraman.
- ▶ *"The good news is that Big Data is here. The bad news is that we are struggling to store and analyze it."* Tom White.

Check out <http://en.wikipedia.org/wiki/Petabyte>


Big Data

Big Data knows everything



Big Data

Friday August 19, 2016 *Boss Freestyles With Jargon*



Big Data



Word-count: Hello World of Big Data



Problem: Given a collection of text files, find the frequency of each word.

Word-count: Hello World of Big Data



Problem: Given a collection of text files, find the frequency of each word.

For example:

File1.txt

File2.txt

File3.txt

large

data

huge

big

data

big

large

data

small

big

data

deluge

Word-count: Hello World of Big Data



Problem: Given a collection of text files, find the frequency of each word.

For example:

File1.txt	File2.txt	File3.txt
large	data	huge
big	data	big
large	data	small
big	data	deluge

Result:

```
large 2
big 3
data 4
small 1
deluge 1
huge 1
```

Questions:

Word-count: Hello World of Big Data



Problem: Given a collection of text files, find the frequency of each word.

For example:

File1.txt	File2.txt	File3.txt
large	data	huge
big	data	big
large	data	small
big	data	deluge

Result:

```
large 2
big 3
data 4
small 1
deluge 1
huge 1
```

Questions: Do we want the output sorted by frequency?

Word-count: Hello World of Big Data



Problem: Given a collection of text files, find the frequency of each word.

For example:

File1.txt	File2.txt	File3.txt
large	data	huge
big	data	big
large	data	small
big	data	deluge

Result:

```
large 2
big 3
data 4
small 1
deluge 1
huge 1
```

Questions: Do we want the output sorted by frequency? Sorted by word?

Word-count: Hello World of Big Data



Problem: Given a collection of text files, find the frequency of each word.

For example:

File1.txt	File2.txt	File3.txt
large	data	huge
big	data	big
large	data	small
big	data	deluge

Result:

```
large 2
big 3
data 4
small 1
deluge 1
huge 1
```

Questions: Do we want the output sorted by frequency? Sorted by word?
How would you solve this problem?

Large Scale Word-Count (1)

- ▶ What if the the number of files is in millions and will not fit in one server?

Large Scale Word-Count (1)

- ▶ What if the the number of files is in millions and will not fit in one server?
- ▶ What if the total size of the files is in Petabytes and will not fit in one server?

Large Scale Word-Count (1)

- ▶ What if the the number of files is in millions and will not fit in one server?
- ▶ What if the total size of the files is in Petabytes and will not fit in one server?
- ▶ How do you modify your solution from before? Assume that you have a cluster of n servers available with the files distributed across the servers.

Large Scale Word-Count (1)

- ▶ What if the the number of files is in millions and will not fit in one server?
- ▶ What if the total size of the files is in Petabytes and will not fit in one server?
- ▶ How do you modify your solution from before? Assume that you have a cluster of n servers available with the files distributed across the servers.
- ▶ But how do we create a cluster and get the files on it?

Large Scale Word-Count (2)

- ▶ What if some of the servers fail while running your program?

Large Scale Word-Count (2)

- ▶ What if some of the servers fail while running your program?
- ▶ What if some of the server disks fail or get corrupted while your program is running?

Large Scale Word-Count (2)

- ▶ What if some of the servers fail while running your program?
- ▶ What if some of the server disks fail or get corrupted while your program is running?
- ▶ What if the some system administrator reboots some of your servers for software/hardware updates without letting you know?

Data Scientist versus Data Engineer

Data Science tasks:

- ▶ Goal is to answer a question or discovering insights.

Data Scientist versus Data Engineer

Data Science tasks:

- ▶ Goal is to answer a question or discovering insights.
- ▶ Often uses interactive shells for adhoc analysis

Data Scientist versus Data Engineer

Data Science tasks:

- ▶ Goal is to answer a question or discovering insights.
- ▶ Often uses interactive shells for adhoc analysis
- ▶ Typically uses Python, R, Matlab (and Spark now)

Data Scientist versus Data Engineer

Data Science tasks:

- ▶ Goal is to answer a question or discovering insights.
- ▶ Often uses interactive shells for adhoc analysis
- ▶ Typically uses Python, R, Matlab (and Spark now)

Data engineer tasks:

- ▶ Builds and maintains a production application (that may use hardened versions of the original data science work)

Data Scientist versus Data Engineer

Data Science tasks:

- ▶ Goal is to answer a question or discovering insights.
- ▶ Often uses interactive shells for adhoc analysis
- ▶ Typically uses Python, R, Matlab (and Spark now)

Data engineer tasks:

- ▶ Builds and maintains a production application (that may use hardened versions of the original data science work)
- ▶ Use principles of software engineering like encapsulation, object-oriented design and interface design

Data Scientist versus Data Engineer

Data Science tasks:

- ▶ Goal is to answer a question or discovering insights.
- ▶ Often uses interactive shells for adhoc analysis
- ▶ Typically uses Python, R, Matlab (and Spark now)

Data engineer tasks:

- ▶ Builds and maintains a production application (that may use hardened versions of the original data science work)
- ▶ Use principles of software engineering like encapsulation, object-oriented design and interface design
- ▶ They have to deal with parallelization, complexity of distributed systems, fault tolerance etc

Data Scientist versus Data Engineer

Data Science tasks:

- ▶ Goal is to answer a question or discovering insights.
- ▶ Often uses interactive shells for adhoc analysis
- ▶ Typically uses Python, R, Matlab (and Spark now)

Data engineer tasks:

- ▶ Builds and maintains a production application (that may use hardened versions of the original data science work)
- ▶ Use principles of software engineering like encapsulation, object-oriented design and interface design
- ▶ They have to deal with parallelization, complexity of distributed systems, fault tolerance etc
- ▶ Typical languages would be Java (with Hadoop and/or Spark), Scala, Python (at smaller scales)