

MapReduce for Large Scale Computing Solutions to Think-Pair-Share Activities



1. MapReduce Exercise 1:

Case Analysis or Capitalization Probability: In a collection of text documents, find the percentage capitalization for each letter of the alphabet. That is, (number of occurrences of a letter that are capitalized/ total number of occurrences for that letter) $\times 100$.

```
file1: The happy Fox
file2: The THE THE
file3: And AND AND
```

```
result:
a: 3/4 * 100 = 75%
t: 4/4 * 100 = 100%
e: 1/4 * 100 = 25%
. . .
```

Solution

The main idea is to output a 2-tuple for both a lower case and upper case character that has the same key: the character in uppercase. The uppercase 2-tuple will have a value of "1" where as the lowercase one would have a value of "0". Then all the 2-tuples for a character will go to one reducer that can add the number of characters as well the number of upper case characters to finally output the percentage.

```
map(String key, String value):
// key: document name
// value: document contents (line by line)
1. for each character ch in value:
2. if (ch is uppercase)
3.     emitIntermediate(ch, 1)
4. else
5.     emitIntermediate(toUppercase(ch), 0)

reduce(String key, Iterable values):
// key: a character
// values: a list of counts
1. int total = 0
2. int sum = 0;
3. for v in values:
4.     sum += parseInt(v)
5.     total += 1
6. emit(key, asString(count * 100.0/total))
```

2. Top-N patents

- Find the number of citations for each patent in a patent reference data set. The format of the input is:

`citing_patent, cited_patent`

Describe a MapReduce algorithm to solve this problem.

```
map(String key, String value):  
// key: document name  
// value: citing_patent, cited_patent (on a line by itself)  
1. parse into two variables: cited_patent and citing_patent  
2. emitIntermediate(cited_patent, 1)  
  
reduce(String key, Iterable values):  
// key: cited_patent  
// values: a list of 1's  
1. int sum = 0;  
2. for v in values:  
3.     sum += parseInt(v)  
4. emit(key, asString(sum))
```

- Find the top N most frequently cited patents. The format of the input is:

`citing_patent, cited_patent`

Describe a MapReduce algorithm to solve this problem.

Hint: This will take two passes.

Solution.

```
map1(String key, String value):
// key: document name
// value: citing_patent, cited_patent (on a line by itself)
1. parse into two variables: cited_patent and citing_patent
2. emitIntermediate(cited_patent, 1)

setup(): //for reduce1
1. initialize a global priority queue for (patent, count) values
   , sorted by count

reduce1(String key, Iterable values):
// key: cited_patent
// values: a list of 1's
1. int sum = 0;
2. for v in values:
3.     sum += parseInt(v)
4. // the priority queue is sorted by sum and only has N
   entries
5. insert (key, sum) into a global priority queue of fixed size
   (N)

cleanup(): //for reduce1
1. for each (patent, count) pair in the global priority queue:
2.     emit(patent, count)

map2(String key, String value):
// key: document name
// value: cited_patent, count
1. parse into two variables: cited_patent and count
2. emitIntermediate(cited_patent, count)

reduce2 --> same as reduce1 (including the same setup and cleanup
)
```