MapReduce for Large Scale Computing Solutions to Think-Pair-Share Activities

1. **MapReduce Exercise 1:**

   Case Analysis or Capitalization Probability: In a collection of text documents, find the percentage capitalization for each letter of the alphabet. That is, (number of occurrences of a letter that are capitalized/ total number of occurrences for that letter) $\times$ 100.

   ```
   file1: The happy Fox
   file2: The THe THE
   file3: And ANd AND

   result:
   a: 3/4 * 100 = 75%
   t: 4/4 * 100 = 100%
   e: 1/4 * 100 = 25%
   . . .
   ```

   Solution

   The main idea is to output a 2-tuple for both a lower case and upper case character that has the same key: the character in uppercase. The uppercase 2-tuple will have a value of "1" where as the lowercase one would have a value of "0". Then all the 2-tuples for a character will go to one reducer that can add the number of characters as well the number of upper case characters to finally output the percentage.

   ```
   map(String key, String value):
   // key: document name
   // value: document contents (line by line)
   1.  for each character ch in value:
   2.  if (ch is uppercase)
   3.      emitIntermediate(ch, 1)
   4.  else
   5.      emitIntermediate(toUppercase(ch), 0)


   reduce(String key, Iterable values):
   // key: a character
   // values: a list of counts
   1.  int total = 0
   2.  int sum = 0;
   3.  for v in values:
   4.      sum += parseInt(v)
   5.      total += 1
   6.  emit(key, asString(count * 100.0/total))
   ```

2. **Top-N patents**

- Find the number of citations for each patent in a patent reference data set. The format of the input is
  citing_patent, cited_patent
  Describe a MapReduce algorithm to solve this problem.

```
map(String key, String value):
// key: document name
// value: citing\_patent, cited\_patent (on a line by itself)
1.   parse into two variables: cited\_patent and citing\_patent
2.   emitIntermediate(cited\_patent, 1)


reduce(String key, Iterable values):
// key: cited\_patent
// values: a list of 1's
1.    int sum = 0;
2.    for v in values:
3.        sum += parseInt(v)
4.    emit(key, asString(sum))
```

- Find the top *N* most frequently cited patents. The format of the input is
  citing_patent, cited_patent
  Describe a MapReduce algorithm to solve this problem.
  Hint: This will take two passes.

  Solution.

```
map1(String key, String value):
// key: document name
// value: citing_patent, cited_patent (on a line by itself)
1.   parse into two variables: cited_patent and citing_patent
2.   emitIntermediate(cited_patent, 1)


reduce1(String key, Iterable values):
// key: cited_patent
// values: a list of 1's
1.    int sum = 0;
2.    for v in values:
3.        sum += parseInt(v)
4.    insert (key, sum) into a global priority queue of fixed size (N)
5.    // the priority queue is sorted by sum and only has N entries

setup(): //for reduce1
1.   initialize a global priority queue for (patent, count) values, sorted by count

cleanup(): //for reduce1
1.   for each (patent, count) pair in the global priority queue:
2.        emit(patent, count)

map2(String key, String value):
// key: document name
// value: cited_patent, count
1.   parse into two variables: cited_patent and count
2.   emitIntermediate(cited_patent, count)

reduce2 --> same as reduce1 (including setup and cleanup)
```