

Spark on Clusters

Spark Cluster Options

- ▶ **Standalone** – a simple cluster manager included with Spark that makes it easy to set up a cluster.
- ▶ **Apache Mesos** – a general cluster manager that can also run Hadoop MapReduce and service applications.
- ▶ **Hadoop YARN** – the resource manager in Hadoop.
- ▶ **Kubernetes** – an open-source system for automating deployment, scaling, and management of containerized applications.

Spark on the cscluster

- ▶ Spark is already installed and configured on the master and all of the worker nodes.
- ▶ Once you tunnel in via SSH (for port 8080), the spark master's web UI is at <http://localhost:8080>, where we can also find the master's URL to use for your spark programs.
- ▶ The spark master's URL is <spark://cscluster00.boisestate.edu:7077>. You can pass it as the "master" argument to `SparkContext` or to `spark-submit`.

Now we can run a spark program on the cluster as follows (we are redirecting standard stream 1) – otherwise all of the words will come out on the console!). Note that we are assuming a shared local filesystem so all Spark nodes can access the data.

```
spark-submit --class "WordCount" --master spark://cscluster00.boisestate.edu:7077 wordcount-basic.jar input 2> log 1> output
```

Spark on the cscluster (2)

- ▶ A spark program running on local host can connect with Hadoop HDFS running in pseudo-distributed mode on the same local host.
- ▶ A spark program running on a standalone cluster needs to connect with Hadoop HDFS running on a cluster, typically the same cluster.

```
spark-submit --class "WordCount" --master spark://cscluster00.boisestate.edu:7077 wordcount-hdfs-save.jar hdfs://cscluster00:9000/user/amit/input
```

- ▶ Note that the wordcount-hdfs-save.jar is a modified version of the word count program that gets the input from HDFS using the following:

```
sc.textFile("hdfs://cscluster00.boisestate.edu:9000/user/amit/input");
```

- ▶ Note that the wordcount-hdfs-save.jar is a modified version of the word count program that saves the output to HDFS using the following:

```
counts.saveAsTextFile("hdfs://cscluster00.boisestate.edu:9000/user/amit/output");
```

Using Jupyter Lab remotely on cscluster

- ▶ Use SSH to tunnel to cscluster00 (use your correct Boise State user name if it is different on your machine) as follows. I chose the port 9088 arbitrarily to not conflict with other users. Choose a different port if you get a conflict.

```
ssh -L 9088:localhost:9088 cscluster00
```

- ▶ Then on cscluster00, start jupyter lab as follows to route the interface back to your machine:

```
jupyter lab --no-browser --port=9088
```

- ▶ Then open up a tab in your browser on your local machine and use the following URL (where the token is shown in the output from jupyter lab when you start it as above.

```
localhost:9088/lab?token=...
```

Standalone Cluster (1)

- ▶ Spark needs to be installed on the master and all of the nodes. For the lab, we installed Spark on `cscluster00`. Since all nodes share the home file system, there is no need to install Spark on the nodes.
- ▶ Start the master (typically, we would use `cscluster00.boisestate.edu`)

```
[amit@cscluster00 ~]$ start-master.sh
starting org.apache.spark.deploy.master.Master, logging to
/home/amit/spark-install/spark/logs/spark-amit-org.apache.spark.deploy.
master.Master-1-cscluster00.boisestate.edu.out
```

- ▶ Once started, the master will print out a `spark://HOST:PORT` URL for itself, which you can use to connect workers to it, or pass as the “master” argument to `SparkContext` or to `spark-submit`.

```
[amit@cscluster00 ~]$ cat /home/amit/spark-install/spark/logs/spark-amit-
-org.apache.spark.deploy.master.Master-1-cscluster00.boisestate.edu
.out | grep spark:
19/11/13 23:48:56 INFO Master: Starting Spark master at spark://
cscluster00.boisestate.edu:7077
```

- ▶ The master’s web UI is at `http://localhost:8080`, where we can also find the master’s URL.

Standalone Cluster (2)

- ▶ Start one or more workers on multiple nodes by passing in the Spark URL of the master. We need password-less ssh access.

```
[amit@cscluster00 ~]$ ssh cscluster01
```

```
Last login: Tue Nov 12 14:27:47 2019 from cscluster00.boisestate.edu
```

```
[amit@cscluster01 ~]$ start-worker.sh spark://cscluster00.boisestate.edu:7077
```

```
starting org.apache.spark.deploy.worker.Worker, logging to /home/amit/spark-install/spark/logs/spark-amit-org.apache.spark.deploy.worker.Worker-1-cscluster01.boisestate.edu.out
```

- ▶ Once you have started a worker, look at the master's web UI. You should see the new node listed there, along with its number of CPUs and memory (minus one gigabyte left for the OS).
- ▶ We can now run Spark jobs on the cluster. For example:

```
spark-submit --class "WordCount" --master spark://cscluster00.boisestate.edu:7077 wordcount-basic.jar input 2> log 1> output
```

- ▶ Finally, we can stop the master and worker daemons.

```
[amit@cscluster00 ~]$ ssh cscluster01 stop-worker.sh
```

```
stopping org.apache.spark.deploy.worker.Worker
```

```
[amit@cscluster00 ~]$ stop-master.sh
```

```
stopping org.apache.spark.deploy.master.Master
```

Standalone Cluster (3)

- We can automate the start/stop of a standalone cluster. Add the names of the worker nodes to the file `~/spark-install/spark/conf/workers`. Then use the `start-all.sh` and `stop-all.sh` scripts.

```
[amit@cscluster00 ~]$ start-all.sh
starting org.apache.spark.deploy.master.Master, logging to /home/amit/
spark-install/spark/logs/spark-amit-org.apache.spark.deploy.master.
Master-1-cscluster00.boisestate.edu.out
cscluster03.boisestate.edu: starting org.apache.spark.deploy.worker.
Worker, logging to /home/amit/spark-install/spark/logs/spark-amit-
org.apache.spark.deploy.worker.Worker-1-cscluster03.boisestate.edu.
out
cscluster01.boisestate.edu: starting org.apache.spark.deploy.worker.
Worker, logging to /home/amit/spark-install/spark/logs/spark-amit-
org.apache.spark.deploy.worker.Worker-1-cscluster01.boisestate.edu.
out
```

```
[amit@cscluster00 ~]$ stop-all.sh
cscluster03.boisestate.edu: stopping org.apache.spark.deploy.worker.
Worker
cscluster01.boisestate.edu: stopping org.apache.spark.deploy.worker.
Worker
stopping org.apache.spark.deploy.master.Master
```


Standalone Cluster with HDFS

- ▶ A spark program running on local host can connect with Hadoop HDFS running in pseudo-distributed mode on the same local host.
- ▶ A spark program running on a standalone cluster needs to connect with Hadoop HDFS running on a cluster, typically the same cluster.
- ▶ Now run the Spark program that accesses the HDFS. The submit command will stay the same.