

Spark on Clusters

Spark Cluster Options

- ▶ **Standalone** – a simple cluster manager included with Spark that makes it easy to set up a cluster.
- ▶ **Apache Mesos** – a general cluster manager that can also run Hadoop MapReduce and service applications.
- ▶ **Hadoop YARN** – the resource manager in Hadoop.
- ▶ **Kubernetes** – an open-source system for automating deployment, scaling, and management of containerized applications.

Spark on the cscluster

- ▶ Spark is already installed and configured on the master and all of the worker nodes.
- ▶ Once we tunnel in via SSH (for port 8080), the Spark master's web UI is at <http://localhost:8080>, where we can also find the master's URL to use for our Spark programs.

```
ssh -L 8080:localhost:8080 cscluster00
```

- ▶ The Spark master's URL is <spark://cscluster00.boisestate.edu:7077>. We can pass it as the “master” argument to `SparkContext` or to `spark-submit`.
- ▶ Now we can run a Spark program on the cluster as follows (we are redirecting standard stream 1 – otherwise all of the words will come out on the console!). Note that we are assuming a shared local filesystem so all Spark nodes can access the data.

```
spark-submit --class "WordCount" --master spark://cscluster00.  
boisestate.edu:7077 wordcount-basic.jar input 2> log 1>  
output
```

Spark on the cscluster (2)

- ▶ A Spark program running on local host can connect with Hadoop HDFS running in pseudo-distributed mode on the same local host.
- ▶ A Spark program running on a standalone cluster needs to connect with Hadoop HDFS running on a cluster, typically the same cluster.

```
spark-submit --class "WordCount" --master spark://cscluster00.  
boisestate.edu:7077 wordcount-hdfs-save.jar hdfs://  
cscluster00:9000/user/amit/input
```

- ▶ Note that the wordcount-hdfs-save.jar is a modified version of the word count program that gets the input from HDFS using the following:

```
sc.textFile("hdfs://cscluster00.boisestate.edu:9000/user/amit/  
input");
```

- ▶ Note that the wordcount-hdfs-save.jar is a modified version of the word count program that saves the output to HDFS using the following:

```
counts.saveAsTextFile("hdfs://cscluster00.boisestate.edu:9000/  
user/amit/output");
```

Using Jupyter Lab remotely on cscluster

- ▶ Python (version 3,8), Jupyter lab, PySpark (version 3.4.1), findspark are all installed on the cluster so we don't have to install anything.
- ▶ Use SSH to tunnel to cscluster00 (we will use our correct Boise State user name if it is different on our machine) as follows. I chose the port 9088 arbitrarily so as to not conflict with other users. Choose a different port if you get a conflict.

```
ssh -L 9088:localhost:9088 cscluster00
```

- ▶ Then on cscluster00, start jupyter lab as follows to route the interface back to our machine:

```
jupyter lab --no-browser --port=9088
```

- ▶ Then open up a tab in a browser on our local machine and use the following URL (where the token is shown in the output from jupyter lab when we start it as above.

```
localhost:9088/lab?token=...
```

Installing Spark on a Standalone Cluster (1)

- ▶ Spark needs to be installed on the master and all of the nodes. For the lab, we installed Spark on `cscluster00`. Since all nodes share the home file system, there is no need to install Spark on the nodes.
- ▶ Start the master (typically, we would use `cscluster00.boisestate.edu`)

```
[amit@cscluster00 ~]$ start-master.sh
starting org.apache.spark.deploy.master.Master, logging to
/home/amit/spark-install/spark/logs/spark-amit-org.apache.spark.deploy.
master.Master-1-cscluster00.boisestate.edu.out
```

- ▶ Once started, the master will print out a `spark://HOST:PORT` URL for itself, which we can use to connect workers to it, or pass as the “master” argument to `SparkContext` or to `spark-submit`.

```
[amit@cscluster00 ~]$ cat /home/amit/spark-install/spark/logs/spark-amit
-org.apache.spark.deploy.master.Master-1-cscluster00.boisestate.edu
.out | grep spark:
19/11/13 23:48:56 INFO Master: Starting Spark master at spark://
cscluster00.boisestate.edu:7077
```

- ▶ The master’s web UI is at `http://localhost:8080`, where we can also find the master’s URL.

Installing Spark on a Standalone Cluster (2)

- ▶ Start one or more workers on multiple nodes by passing in the Spark URL of the master. We need password-less ssh access.

```
[amit@cscluster00 ~]$ ssh cscluster01
```

```
Last login: Tue Nov 12 14:27:47 2019 from cscluster00.boisestate.edu
```

```
[amit@cscluster01 ~]$ start-worker.sh spark://cscluster00.boisestate.edu:7077
```

```
starting org.apache.spark.deploy.worker.Worker, logging to /home/amit/spark-install/spark/logs/spark-amit-org.apache.spark.deploy.worker.Worker-1-cscluster01.boisestate.edu.out
```

- ▶ Once we have started a worker, look at the master's web UI. We should see the new node listed there, along with its number of CPUs and memory (minus one gigabyte left for the OS).
- ▶ We can now run Spark jobs on the cluster. For example:

```
spark-submit --class "WordCount" --master spark://cscluster00.boisestate.edu:7077 wordcount-basic.jar input 2> log 1> output
```

- ▶ Finally, we can stop the master and worker daemons.

```
[amit@cscluster00 ~]$ ssh cscluster01 stop-worker.sh
```

```
stopping org.apache.spark.deploy.worker.Worker
```

```
[amit@cscluster00 ~]$ stop-master.sh
```

```
stopping org.apache.spark.deploy.master.Master
```

Installing Spark on a Standalone Cluster (3)

- We can automate the start/stop of a standalone cluster. Add the names of the worker nodes to the file `~/spark-install/spark/conf/workers`. Then use the `start-all.sh` and `stop-all.sh` scripts.

```
[amit@cscluster00 ~]$ start-all.sh
starting org.apache.spark.deploy.master.Master, logging to /home/amit/
spark-install/spark/logs/spark-amit-org.apache.spark.deploy.master.
Master-1-cscluster00.boisestate.edu.out
cscluster03.boisestate.edu: starting org.apache.spark.deploy.worker.
Worker, logging to /home/amit/spark-install/spark/logs/spark-amit-
org.apache.spark.deploy.worker.Worker-1-cscluster03.boisestate.edu.
out
cscluster01.boisestate.edu: starting org.apache.spark.deploy.worker.
Worker, logging to /home/amit/spark-install/spark/logs/spark-amit-
org.apache.spark.deploy.worker.Worker-1-cscluster01.boisestate.edu.
out
```

```
[amit@cscluster00 ~]$ stop-all.sh
cscluster03.boisestate.edu: stopping org.apache.spark.deploy.worker.
Worker
cscluster01.boisestate.edu: stopping org.apache.spark.deploy.worker.
Worker
stopping org.apache.spark.deploy.master.Master
```


Running Spark using Hadoop YARN scheduler

- ▶ We can redirect Spark to use Hadoop YARN as the job scheduler by first setting the following environment variable.

```
export HADOOP_CONF_DIR=$HADOOP_HOME/etc/hadoop
```

- ▶ Then we can simply change the `--master` option for `spark-submit` to be `yarn` and viola! For example:

```
spark-submit --master yarn --class WordCount wordcount-hdfs-  
save.jar hdfs://cscluster00:9000/user/amit/input
```

- ▶ To make Hadoop YARN the default for Spark, we can do the following steps to set the default.

```
mv $SPARK_HOME/conf/spark-defaults.conf.template $SPARK_HOME  
/conf/spark-defaults.conf  
//edit spark-defaults.conf and add the following line  
spark-master yarn
```

Running Spark on a Cluster from Our Laptop!

- ▶ Spark jobs can run on YARN in two modes: **cluster** mode and **client** mode.
- ▶ Recall that a Spark job consists of two parts: Spark Executors that run the actual tasks, and a Spark Driver that schedules the Executors.
 - ▶ **cluster** mode: everything runs inside the cluster. We can start a job from our laptop and the job will continue running even if we close our computer.
 - ▶ **client** mode: the Spark driver runs on a client, such as our laptop. If the client is shut down, the job fails.
- ▶ Client mode is well suited for interactive jobs, but applications will fail if the client stops. For long running jobs, cluster mode is more appropriate.

Examples of Remote Submission to the Cluster

- ▶ Spark driver running as a client on our laptop, the executors are running on the Spark cluster. Note that the deploy mode is `client` by default.

```
spark-submit --class "WordCount" --deploy-mode client --master spark://  
cscluster00.boisestate.edu:7077 wordcount-hdfs-save.jar hdfs://  
cscluster00:9000/user/amit/input
```

- ▶ Spark driver running on the master, the executors running on the Spark cluster. Note that the our laptop is just being used to submit the job. It will not see any standard output/error messages from the driver. Note also that we have to specify where to find the jar file on the cluster and not locally.

```
spark-submit --class "WordCount" --deploy-mode cluster --master spark://  
cscluster00.boisestate.edu:7077 ~/CS535-resources/examples/spark/  
word-count/wordcount-hdfs-save.jar hdfs://cscluster00:9000/user/  
amit/input
```

- ▶ In order to remotely submit Spark jobs to Hadoop YARN, we will need to copy the hadoop config from the cluster to a local folder and then set `HADOOP_CONF_DIR` to point to this folder. Then we can use `--master yarn` option. Note that this will break our local pseudo-distributed Hadoop setup until we reset the environment variable.