

Spark SQL

Introduction

- ▶ **Spark SQL** is a module for structured data processing. It provides Spark with more information about the both the data and the computation being performed.
- ▶ We can interact with Spark SQL with *SQL*, *Dataset* (and *DataFrames*) API and other ways.
 - ▶ We can interact with the SQL interface via a programming language and the data set will be returned as a *Dataset/DataFrame*. We can also interact with the SQL interface using the *command-line* or over *JDBC/ODBC*.
- ▶ When computing a result the same execution engine is used, independent of which API/language you are using to express the computation. This unification means that developers can easily switch back and forth.

Creating DataFrames

- ▶ With a [SparkSession](#), we can create [DataFrames](#) from an existing RDD, from a Hive table or from other Spark data sources.
- ▶ [JavaSparkSQLException1.java](#) in the folder Spark/java-sql-example
- ▶ [DataFrames](#) provide a domain-specific language for structured data manipulation. [DataFrames](#) are just [Dataset](#) of Rows in Scala and Java API. These operations are also referred as “untyped transformations” in contrast to “typed transformations” that come with strongly typed Scala/Java [Datasets](#).
- ▶ See same example as above for some basic examples of structured data processing using [Datasets](#). Here are some code snippets:

```
df.printSchema();
```

```
// Select only the "name" column
df.select("name").show();
```

```
// Select everybody, but increment the age by 1
df.select(col("name"), col("age").plus(1)).show();
```

```
// Select people older than 21
df.filter(col("age").gt(21)).show();
```

```
// Count people by age
df.groupBy("age").count().show();
```

Running SQL Queries Programmatically

- ▶ The `sql` function on a `SparkSession` enables applications to run SQL queries programmatically and returns the result as a `Dataset<Row>`. See same example as before. Here is a code snippet.

```
df.createOrReplaceTempView("people");  
Dataset<Row> sqlDF = spark.sql("SELECT * FROM people"  
    );  
sqlDF.show();
```

- ▶ Create a global temporary view to have one that is shared among all sessions while a Spark application runs. Temporary views (like above) are session-scoped and will disappear if the session that creates it terminates.

```
df.createGlobalTempView("people");  
spark.sql("SELECT * FROM global_temp.people").show();  
spark.newSession().sql("SELECT * FROM global_temp.  
    people").show();
```

Creating **Datasets**

- ▶ **Datasets** are similar to RDDs, however, instead of using Java serialization or Kryo they use a specialized **Encoder** to serialize the objects for processing or transmitting over the network
- ▶ Encoders are code generated dynamically and use a format that allows Spark to perform many operations like filtering, sorting and hashing without deserializing the bytes back into an object.
- ▶ See example: [JavaSQLExample2.java](#)

Interoperability with RDDs (1)

- ▶ Spark SQL supports two different methods for converting existing RDDs into **Datasets**.
 - ▶ Use reflection to infer the schema of an RDD that contains specific types of objects. This reflection-based approach leads to more concise code and works well when we already know the schema.
 - ▶ Use a programmatic interface that allows us to construct a schema and then apply it to an existing RDD. While this method is more verbose, it allows us to construct **Datasets** when the columns and their types are not known until run time.

Interoperability with RDDs (2)

- ▶ Spark SQL supports automatically converting an RDD of `JavaBeans` into a `DataFrame`.
- ▶ The `BeanInfo`, obtained using reflection, defines the schema of the table. Currently, Spark SQL does not support `JavaBeans` that contain Map field(s). Nested `JavaBeans` and `List` or `Array` fields are supported though.
- ▶ See the method `runInferSchemaExample` in `JavaSQLExample3.java`

What is a `JavaBean`?

- ▶ It is a reusable software component written in Java.
- ▶ `JavaBeans` are classes that encapsulate many objects into a single object (the bean) by following a standard. The name "Bean" was given to encompass this standard, which is defined below:
 - ▶ Must implement `Serializable`
 - ▶ It should have a public no-arg constructor
 - ▶ All attributes must be private with public `getters` and `setter` methods
 - ▶ Rules for `setter` methods:
 - ▶ It should be public
 - ▶ The return-type should be void
 - ▶ The method name should be prefixed with set
 - ▶ It should take some argument i.e. it should not be no-arg method
 - ▶ Rules for `getter` methods:
 - ▶ It should be public
 - ▶ The return-type should not be void
 - ▶ The method name should be prefixed with get
 - ▶ For boolean attributes, the name can be prefixed with "get" or "is" but preferred prefix would be "is"
 - ▶ It should not take any argument

Interoperability with RDDs (3)

- ▶ When `JavaBean` classes cannot be defined ahead of time, a `Dataset<Row>` can be created programmatically with three steps.
 - ▶ **Step 1:** Create an RDD of Rows from the original RDD
 - ▶ **Step 2:** Create the schema represented by a `StructType` matching the structure of Rows in the RDD created in Step 1
 - ▶ **Step 3:** Apply the schema to the RDD of Rows via `createDataFrame` method provided by `SparkSession`
- ▶ See the method `runProgrammaticSchemaExample` in `JavaSQLExample3.java`